

Temporal databases

Storing the history of data


Learning goals

1. Understand the need for modelling temporal changes in a database
2. Recognise the three axes of time
3. Learn to create a temporal database using the tools provided by a RDMS.

Changes in time

- Earlier, you have learnt logical database design, i.e. the design of the database schema.
- In real world, things happen in time, and as a consequence, data changes.
- As data changes, historical data accumulates.
 - By historical data, we mean old (updated or deleted) values of the data items.
 - In some contexts, it may be vital to store the historical data.
- Possible queries to historical data:
 - What earlier values has the data item had?
 - At what time has each change of the data item taken place?
 - What has been the period of validity for each data point in history?
- Most real-life databases contain a temporal dimension at least to some extent.
- Its modelling by tables and fields in the schema may be challenging.
- The database management systems (such as MariaDB) provide tools to temporal modelling.

A database without a temporal dimension


id		firstname	lastname	salary
1		Viivi	Virta	2 000,00
2		Ahmed	Bakir	2 000,00

- In the example, there is a simple Employee table.
- Let's assume that Viivi gets a pay rise to 2200 euros.
- Simple solution: update the value of the salary field.
- The history of Viivi's salary is not saved. For example, the following challenges arise:
 1. What to do if the history of her salary needs to be investigated later?
 - For example, how to retrieve a salary that was valid on 14 March 2023?
 2. What to do when Viivi's salary is to be raised to 2500 euros after two months, as her probation ends?
 - Can the new salary somehow be updated in advance?

Temporal database


- Problems such as those described before are avoided by creating a temporal database
 - temporal = time-related.
- Temporal databases may include three axes of time:
 1. transaction time
 2. valid time
 3. decision time
- Of these, one or more axis can be applied simultaneously. Often, the transaction time and the valid time, or either of them, is used.

Axis 1/3: transaction time

id		firstname	lastname	salary	transaction_time
1		Viivi	Virta	2 000,00	2022-08-12 10:14:41.649413
1		Viivi	Virta	2 300,00	2022-08-12 10:14:44.708788
2		Ahmed	Bakir	2 000,00	2022-08-12 10:14:41.654144


- The transaction time tells
 - The point in time when the record has been stored in the database, or
 - The period of time when the record has been current (i.e. the last record in the table)
- In the example, Viivi's salary has been updated three seconds after the initial insert.
- The table now includes the transaction time axis (of a temporal database)
- Historical values as well as update times will be stored.
- The transaction time is not necessarily the time when the updated data has been valid from an application's point of view.
 - Viivi may have started her job in the beginning of August. Her salary may well be 2300 €/mo. from the beginning of August.
- Is recording the transaction time a valid solution to store the salary history?
 - It may be, if all updates are made between consecutive paydays, and the new value is always valid for the entire month.
 - However, future payraises cannot be stored e.g. two months in advance.

Axis 2/3: valid time

id		firstname	lastname	salary	startdate	enddate
1		Viivi	Virta	2 500,00	2022-11-01	2024-12-31
2		Ahmed	Bakir	2 000,00	2022-08-01	2024-12-31
3		Viivi	Virta	2 000,00	2022-08-01	2022-11-01

- Valid time, or application time, tells when the data is valid from the point of view of data usage.
- From this point of view, it is irrelevant when the data has been stored or updated in the database.

Axis 3/3: decision time

id		firstname	lastname	salary	decision_time
1		Viivi	Virta	2 000,00	2022-07-24
2		Ahmed	Bakir	2 000,00	2022-07-31

- The decision time tells when the decision about the validity of the record has been made.
 - It can differ from both the transaction time and the decision time.
- An example of the three axes of time:
- In the board meeting on 24.7.2022 it was decided that Viivi's salary will be 2000 €/mo. in the time period of 1.8.2022 to 31.12.2022. This fact was entered into the system on 27.7.2022.
 - Decision time: 24.7.2022
 - Valid time: 1.8.2022 to 31.12.2022
 - Transaction time: 27.7.2022
- From the database's point of view, the decision time is stored in a regular date field. Thus, it is not particularly interesting when technical aspects are considered.

SQL and axes of time

- SQL:2011 standard defines operations for a bitemporal database.
 - A bitemporal database contains both the valid time and the transaction time.
- In MariaDB, support for bitemporal databases has been present since version 10.3.4.
 - system versioning
 - application-time period versioning
- More functionality is included as versions evolve: for example, since version 10.5.3, there has been support for WITHOUT OVERLAPS clause that prevents overlapping intervals in the valid time.
 - A potential use case: reservations to a hotel room. Two reservations for the same room must not overlap.
- Next, let's get acquainted with the use of mobth axes of time in MariaDB.

Turning transaction time on

```
CREATE TABLE person (  
    id INT NOT NULL AUTO_INCREMENT,  
    firstname VARCHAR(40),  
    lastname VARCHAR(40),  
    salary DECIMAL(8,2),  
    PRIMARY KEY (id)  
) WITH SYSTEM VERSIONING;
```

- The WITH SYSTEM VERSIONING clause in the end turns on the system versioning based on the transaction time.
- There will be hidden fields row_start and row_end in the table. They are timestamp fields for the transaction time (i.e. the start and end of the period when the record has been current)
 - These fields are invisible by default (SELECT * ... won't display them).

Queries in presence of the transaction time

```
INSERT INTO person VALUES (NULL, "Viivi", "Virta", 2000);  
INSERT INTO person VALUES (NULL, "Ahmed", "Bakir", 2000);
```

1.

```
SELECT * FROM person;  
  
SELECT SLEEP(3);  
UPDATE person SET salary=2300 WHERE id=1;
```
2.

```
SELECT * FROM person;
```
3.

```
SELECT *, row_start AS transaction_time FROM person FOR SYSTEM_TIME ALL;
```
4.

```
SELECT * FROM person FOR SYSTEM_TIME AS OF TIMESTAMP SUBTIME(NOW(), "1.0");
```

1.

id	firstname	lastname	salary
1	Viivi	Virta	2 000,00
2	Ahmed	Bakir	2 000,00

2.

id	firstname	lastname	salary
1	Viivi	Virta	2 300,00
2	Ahmed	Bakir	2 000,00

3.

id	firstname	lastname	salary	transaction_time
1	Viivi	Virta	2 000,00	2022-08-12 12:51:32.725143
1	Viivi	Virta	2 300,00	2022-08-12 12:51:35.780547
2	Ahmed	Bakir	2 000,00	2022-08-12 12:51:32.729443

4.

id	firstname	lastname	salary
1	Viivi	Virta	2 000,00
2	Ahmed	Bakir	2 000,00

Queries in presence of the transaction time

- With system versioning (that is, recording the transaction time), the operations for the current record (SELECT, INSERT, UPDATE, DELETE) can be written as if the transaction time didn't exist.
- A regular SELECT retrieves the current record.
- A regular DELETE marks the record as invalid,. It doesn't delete the record physically from the disk.
 - If needed, all versions can be explicitly deleted with a DELETE HISTORY statement.

Queries based on the transaction time

- It is possible to refer to a historical version by ending a `FOR SYSTEM_TIME` clause in the end of the query.
- Examples:
 - `SELECT ... FOR SYSTEM_TIME ALL` retrieves all historical versions separately.
 - `SELECT ... FOR SYSTEM_TIME SUBTIME(NOW(), "365 0:0:0")` retrieves a version that is 365 days old.
 - `SELECT ... FOR SYSTEM_TIME TIMESTAMP("2022-08-24 18:55")` retrieves a version that has been current on 24.8.2022 at 18:55.

Turning valid time on

```
CREATE TABLE person (  
    id INT NOT NULL AUTO_INCREMENT,  
    firstname VARCHAR(40),  
    lastname VARCHAR(40),  
    salary DECIMAL(8,2),  
    startdate DATE,  
    enddate DATE,  
    PERIOD FOR valid_time(startdate, enddate),  
    PRIMARY KEY (id)  
);
```

- In the table, specific fields will be introduced for both the start and end dates of the validity period.
 - The permitted data types are DATE, DATETIME, and TIMESTAMP. The types must be same for the two fields.
- PERIOD FOR clause turns on the versioning based on the valid time.
 - In the example, valid_time is an identifier with which you can refer to the valid time period later. The naming of the identifier is free.

Queries based on the valid time

```
INSERT INTO person VALUES (NULL, "Viivi", "Virta", 2000, "2022-08-01", "2024-12-31");
INSERT INTO person VALUES (NULL, "Ahmed", "Bakir", 2000, "2022-08-01", "2024-12-31");
```

1. `SELECT * FROM person;`

```
UPDATE person
FOR PORTION OF valid_time FROM "2022-11-01" TO "2024-12-31"
SET salary = 2500
WHERE id=1;
```

Note: in example 3., **NOW()** returns a timestamp in Sep 2022.

2. `SELECT * FROM person;`

3. `SELECT * FROM person WHERE NOW() BETWEEN startdate AND enddate;`

4. `SELECT * FROM person WHERE TIMESTAMP("2022-11-07") BETWEEN startdate AND enddate;`

1.

id	firstname	lastname	salary	startdate	enddate
1	Viivi	Virta	2 000,00	2022-08-01	2024-12-31
2	Ahmed	Bakir	2 000,00	2022-08-01	2024-12-31

2.

id	firstname	lastname	salary	startdate	enddate
1	Viivi	Virta	2 500,00	2022-11-01	2024-12-31
2	Ahmed	Bakir	2 000,00	2022-08-01	2024-12-31
3	Viivi	Virta	2 000,00	2022-08-01	2022-11-01

3.

id	firstname	lastname	salary	startdate	enddate
2	Ahmed	Bakir	2 000,00	2022-08-01	2024-12-31
3	Viivi	Virta	2 000,00	2022-08-01	2022-11-01

4.

id	firstname	lastname	salary	startdate	enddate
1	Viivi	Virta	2 500,00	2022-11-01	2024-12-31
2	Ahmed	Bakir	2 000,00	2022-08-01	2024-12-31

Queries based on the valid time

- SELECT queries are written in a regular fashion. If the records need to be selected on the basis of valid time, the condition is included in the WHERE clause.
- UPDATE and DELETE queries may contain a FOR PORTION clause that changes the data for the given valid given interval only. As a consequence, new versions of the record may emerge:
 - In the previous example, the UPDATE statement created a new period with a higher salary. The old period (with a lower salary) shrank automatically.
 - Note that the UPDATE statement created a new record into the person table that carries a new primary key value. If id needs to be a fixed person number (instead of just a row key without further meaning), it must be made a regular field, not a primary key field.

A bitemporal table

```
DROP TABLE IF EXISTS person;

CREATE TABLE person (
  id INT NOT NULL AUTO_INCREMENT,
  firstname VARCHAR(40),
  lastname VARCHAR(40),
  salary DECIMAL(8,2),
  startdate DATE,
  enddate DATE,
  PERIOD FOR valid_time(startdate, enddate),
  PRIMARY KEY (id)
) WITH SYSTEM VERSIONING;

INSERT INTO person VALUES (NULL, "Viivi", "Virta", 2000, "2022-08-01", "2024-12-31");
INSERT INTO person VALUES (NULL, "Ahmed", "Bakir", 2000, "2022-08-01", "2024-12-31");

SELECT *, row_start, row_end FROM person;
```

- Both axis of time (transaction time and valid time) can be taken into use simultaneously.

id	firstname	lastname	salary	startdate	enddate	row_start	row_end
1	Viivi	Virta	2 000,00	2022-08-01	2024-12-31	2022-08-12 14:05:37.071090	2038-01-19 05:14:07.999999
2	Ahmed	Bakir	2 000,00	2022-08-01	2024-12-31	2022-08-12 14:05:37.077544	2038-01-19 05:14:07.999999

Another option: modelling in the schema

- Before, we discussed the temporal database functionality of a relational database management system (MariaDB).
- Is it possible to achieve the same functionality without the provided functionality?
- Alternatively, the history tables and fields can be modelled in the schema.
 - For each INSERT or UPDATE, the transaction time can be generated with the SQL NOW() function.
- With devotion, the correct schema can be designed. However, the UPDATE and SELECT operations may become cumbersome.

Another option: modelling in the schema

```
CREATE TABLE person (  
    id INT NOT NULL AUTO_INCREMENT,  
    firstname VARCHAR(40),  
    lastname VARCHAR(40),  
    PRIMARY KEY (id)  
);  
  
CREATE TABLE salary (  
    salaryid INT NOT NULL AUTO_INCREMENT,  
    personid INT,  
    transactiontime TIMESTAMP DEFAULT NOW(),  
    salary DECIMAL(8,2),  
    startdate DATE,  
    enddate DATE,  
    PRIMARY KEY (salaryid),  
    FOREIGN KEY (personid) REFERENCES person(id)  
);
```

Another option: modelling in the schema

- The queries and the update operations become more difficult.
- Let's raise Viivi's salary from a date in the future.
 - This operation becomes quite complicated...
 - ... and the things like this become even more difficult as size of the schema grows.

```
UPDATE salary SET enddate="2022-11-01"
WHERE personid=1
AND enddate>"2022-11-01"
AND startdate<="2024-12-31";
INSERT INTO salary(personid, salary, startdate, enddate)
VALUES(1, 2500, "2022-11-01", "2024-12-31");

SELECT * FROM salary;
```



salaryid	personid	transactiontime	salary	startdate	enddate
1	1	2022-08-12 14:49:02	2 000,00	2022-08-01	2022-11-01
2	1	2022-08-12 14:49:05	2 500,00	2022-11-01	2024-12-31

Modelling in the scema: benefits and drawbacks

- Benefits:
 - The temporal nature of data is made explicitly visible in the design phase.
 - The solution is guaranteed to work in any RDBMS.
- Drawbacks:
 - The schema becomes complicated.
 - The queries and the update operations become extremely complex.
 - This complexity is likely to transfer to the application layer.
 - Consequently, the performance may be hampered.

A lightweight solution: logging

- If the queries to the historical data are rare, and not needed in the daily operation, logging may be a sufficient solution.
- In MariaDB, general query logging can be turned on. As a consequence, each SQL statement is written into a text file.
- Historical data is still retrievable, but it is an exceptional procedure that cannot be expressed with normal SQL queries.

Issues to consider

- The need for temporal data should be analyzed in the modelling phase.
- With temporal databases, there is the risk of over-modelling: is history really needed for all potentially changing data.
 - Example: can a name of a Finnish municipality change? Is it necessary to retrieve the old name in its original shape?
- It usually pays to take advantage of the temporal database functionality in a RDBMS.
 - Specifically, MariaDB supports bitemporal tables.